

The Greedy Method



Civil War Knapsack. U.S. government image. Vicksburg National Military Park. Public domain.

The Fractional Knapsack Problem



- ◆ Given: A set S of n items, with each item i having
 - b_i - a positive benefit
 - w_i - a positive weight
- ◆ Goal: Choose items with maximum total benefit but with weight at most W .
- ◆ If we are allowed to take fractional amounts, then this is the **fractional knapsack problem**.
 - In this case, we let x_i denote the amount we take of item i




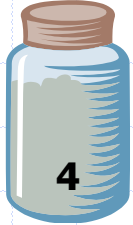

- Objective: maximize
$$\sum_{i \in S} b_i (x_i / w_i)$$

- Constraint:
$$\sum_{i \in S} x_i \leq W$$

Example



- ◆ Given: A set S of n items, with each item i having
 - b_i - a positive benefit
 - w_i - a positive weight
- ◆ Goal: Choose items with maximum total benefit but with weight at most W .

Items:					
Weight:	4 ml	8 ml	2 ml	6 ml	1 ml
Benefit:	\$12	\$32	\$40	\$30	\$50
Value:	3	4	20	5	50
	(\$ per ml)				



“knapsack”

10 ml

Solution:

- 1 ml of 5
- 2 ml of 3
- 6 ml of 4
- 1 ml of 2

The Fractional Knapsack Algorithm



- ◆ Greedy choice: Keep taking item with highest **value** (benefit to weight ratio)
 - Since $\sum_{i \in S} b_i (x_i / w_i) = \sum_{i \in S} (b_i / w_i) x_i$
 - Run time: $O(n \log n)$. Why?
- ◆ Correctness: Suppose there is a better solution
 - there is an item i with higher value than a chosen item j , but $x_i < w_i$, $x_j > 0$ and $v_i < v_j$
 - If we substitute some i with j , we get a better solution
 - How much of i : $\min\{w_i - x_i, x_j\}$
 - Thus, there is no better solution than the greedy one

Algorithm *fractionalKnapsack*(S, W)

Input: set S of items w/ benefit b_i and weight w_i ; max. weight W

Output: amount x_i of each item i to maximize benefit w/ weight at most W

for each item i in S

$x_i \leftarrow 0$

$v_i \leftarrow b_i / w_i$ {value}

$w \leftarrow 0$ {total weight}

while $w < W$

remove item i w/ highest v_i

$x_i \leftarrow \min\{w_i, W - w\}$


$w \leftarrow w + \min\{w_i, W - w\}$

return x


Analysis of Greedy Algorithm for Fractional Knapsack Problem


- ◆ We can sort the items by their benefit-to-weight values, and then process them in this order.
- ◆ This would require $O(n \log n)$ time to sort the items and then $O(n)$ time to process them in the while-loop.
- ◆ To see that our algorithm is correct, suppose, for the sake of contradiction, that there is an optimal solution better than the one chosen by this greedy algorithm.
- ◆ Then there must be two items i and j such that
$$x_i < w_i, x_j > 0, \text{ and } v_i > v_j .$$
- ◆ Let $y = \min\{w_i - x_i, x_j\}$.
- ◆ But then we could replace an amount y of item j with an equal amount of item i , thus increasing the total benefit without changing the total weight, which contradicts the assumption that this non-greedy solution is optimal.

Task Scheduling




- ◆ Given: a set T of n tasks, each having:
 - A start time, s_i
 - A finish time, f_i (where $s_i < f_i$)
- ◆ Goal: Perform all the tasks using a minimum number of “machines.”






Example

- ◆ Given: a set T of n tasks, each having:
 - A start time, s_i
 - A finish time, f_i (where $s_i < f_i$)
 - $[1,4], [1,3], [2,5], [3,7], [4,7], [6,9], [7,8]$ (ordered by start)
- ◆ Goal: Perform all tasks on min. number of machines



Task Scheduling Algorithm



- ◆ Greedy choice: consider tasks by their start time and use as few machines as possible with this order.
 - Run time: $O(n \log n)$. Why?
- ◆ Correctness: Suppose there is a better schedule.
 - We can use $k-1$ machines
 - The algorithm uses k
 - Let i be first task scheduled on machine k
 - Machine i must conflict with $k-1$ other tasks
 - But that means there is no non-conflicting schedule using $k-1$ machines

Algorithm *taskSchedule*(T)

Input: set T of tasks w/ start time s_i and finish time f_i

Output: non-conflicting schedule with minimum number of machines

$m \leftarrow 0$ {no. of machines}

while T is not empty

remove task i w/ smallest s_i

if *there's a machine j for i* **then**

schedule i on machine j

else

$m \leftarrow m + 1$

schedule i on machine m

return *schedule*